

Cloud-Native Data Architecture and Bankruptcy Prediction for MSBA Financial Group

Ritvik Vasikarla: GenBus 780

October 1, 2025

```
[1]: import pyarrow.parquet as pq
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import boto3
from io import StringIO
import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
sns.set_style('whitegrid')

print('Libraries Imported')
```

```
/home/ec2-user/anaconda3/envs/python3/lib/python3.10/site-
packages/pandas/core/computation/expressions.py:21: UserWarning: Pandas requires
version '2.8.4' or newer of 'numexpr' (version '2.7.3' currently installed).
```

```
from pandas.core.computation.check import NUMEXPR_INSTALLED
```

Libraries Imported

```
[2]: s3_client = boto3.client('s3', region_name='us-east-1')
bucket_name = 'msba-financial-data-ritvik'
```

```
[3]: files = s3_client.list_objects_v2(
    Bucket=bucket_name,
    Prefix = 'processed-data/combined_financial_data/train/'
)
```

```
[4]: import io

dfs = []
for obj in files.get('Contents', []):
    if obj['Key'].endswith('.parquet'):
        print(f"    Reading: {obj['Key']}")
        file_obj = s3_client.get_object(Bucket=bucket_name, Key=obj['Key'])
```

```
df_temp = pd.read_parquet(io.BytesIO(file_obj['Body'].read()))
dfs.append(df_temp)
```

```
df = pd.concat(dfs, ignore_index=True)
print('Training Data Loaded')
```

Reading: processed-data/combined_financial_data/train/part-00000-31f68e42-3f3a-4d6c-84f3-0e8dbd0acc64-c000.snappy.parquet

Reading: processed-data/combined_financial_data/train/part-00001-31f68e42-3f3a-4d6c-84f3-0e8dbd0acc64-c000.snappy.parquet

Reading: processed-data/combined_financial_data/train/part-00002-31f68e42-3f3a-4d6c-84f3-0e8dbd0acc64-c000.snappy.parquet

Reading: processed-data/combined_financial_data/train/part-00003-31f68e42-3f3a-4d6c-84f3-0e8dbd0acc64-c000.snappy.parquet

Training Data Loaded

```
[5]: files_test = s3_client.list_objects_v2(
      Bucket=bucket_name,
      Prefix='processed-data/combined_financial_data/test/'
    )
```

```
[6]: dfs_test = []
for obj in files_test.get('Contents', []):
    if obj['Key'].endswith('.parquet'):
        print(f"    Reading: {obj['Key']}")
        file_obj = s3_client.get_object(Bucket=bucket_name, Key=obj['Key'])
        df_temp = pd.read_parquet(io.BytesIO(file_obj['Body'].read()))
        dfs_test.append(df_temp)
df_test = pd.concat(dfs_test, ignore_index=True)
print('Test Data Loaded')
```

Reading: processed-data/combined_financial_data/test/part-00000-a3aa98fc-bf1e-4abe-8ba6-edfe698e4691-c000.snappy.parquet

Reading: processed-data/combined_financial_data/test/part-00001-a3aa98fc-bf1e-4abe-8ba6-edfe698e4691-c000.snappy.parquet

Reading: processed-data/combined_financial_data/test/part-00002-a3aa98fc-bf1e-4abe-8ba6-edfe698e4691-c000.snappy.parquet

Reading: processed-data/combined_financial_data/test/part-00003-a3aa98fc-bf1e-4abe-8ba6-edfe698e4691-c000.snappy.parquet

Test Data Loaded

```
[7]: print(f'Shape: {df.shape}')
      print(f'Shape: {df_test.shape}')
```

Shape: (5522, 24)

Shape: (1297, 24)

```
[8]: bankruptcy_col = 'bankrupt'
      bankruptcy_col
```

```
[8]: 'bankrupt'
```

```
[9]: df.dtypes
```

```
[9]: bankrupt                int32
     company_id             object
     net_worth_to_assets    float64
     retained_earnings_to_total_assets float64
     working_capital_to_total_assets float64
     working_capital_to_equity float64
     equity_to_longterm_liability float64
     current_liabilities_to_equity float64
     liability_to_equity    float64
     current_liability_to_current_assets float64
     borrowing_dependency  float64
     debt_ratio_percentage float64
     persistent_eps        float64
     per_share_net_profit_pre_tax float64
     operating_profit_per_share float64
     tax_rate              float64
     operating_gross_margin float64
     net_income_to_total_assets float64
     roa_before_interest_percent_after_tax float64
     net_profit_before_tax_to_paid_in_capital float64
     net_income_to_stockholders_equity float64
     operating_profit_paid_in_capital float64
     total_asset_turnover float64
     total_expense_to_assets float64
     dtype: object
```

```
[10]: miss = df.isnull().sum()
      if miss.sum() > 0:
          print(miss[miss>0])
      else:
          print('No missing values')
      print()
```

No missing values

```
[11]: df.describe()
```

```
[11]:
```

	bankrupt	net_worth_to_assets	retained_earnings_to_total_assets \
count	5522.000000	5522.000000	5522.000000
mean	0.032235	0.886930	0.934734
std	0.176639	0.054449	0.026537
min	0.000000	0.000000	0.000000
25%	0.000000	0.850772	0.931285

50%	0.000000	0.888729	0.937731
75%	0.000000	0.927888	0.944905
max	1.000000	1.000000	1.000000

	working_capital_to_total_assets	working_capital_to_equity \
count	5522.000000	5522.000000
mean	0.814207	0.735789
std	0.059537	0.012808
min	0.000000	0.000000
25%	0.774267	0.733604
50%	0.810320	0.736003
75%	0.850366	0.738586
max	1.000000	1.000000

	equity_to_longterm_liability	current_liabilities_to_equity \
count	5522.000000	5522.000000
mean	0.115828	0.331506
std	0.021421	0.014028
min	0.025851	0.153811
25%	0.110933	0.328084
50%	0.112375	0.329661
75%	0.117123	0.332314
max	1.000000	1.000000

	liability_to_equity	current_liability_to_current_assets \
count	5522.000000	5522.000000
mean	0.280490	0.031678
std	0.015398	0.032900
min	0.133503	0.000122
25%	0.276913	0.017990
50%	0.278774	0.027595
75%	0.281490	0.038421
max	1.000000	1.000000

	borrowing_dependency	...	operating_profit_per_share	tax_rate \
count	5522.000000	...	5522.000000	5522.000000
mean	0.374813	...	0.109298	0.115100
std	0.017118	...	0.028906	0.139957
min	0.187124	...	0.000000	0.000000
25%	0.370167	...	0.096246	0.000000
50%	0.372585	...	0.104307	0.071971
75%	0.376224	...	0.116196	0.206296
max	1.000000	...	1.000000	1.000000

	operating_gross_margin	net_income_to_total_assets \
count	5522.000000	5522.000000
mean	0.608096	0.807902

std	0.015697	0.040579
min	0.156308	0.000000
25%	0.600457	0.796871
50%	0.606062	0.810916
75%	0.614017	0.826375
max	1.000000	0.982879

	roa_before_interest_percent_after_tax \
count	5522.000000
mean	0.558970
std	0.065884
min	0.000000
25%	0.535979
50%	0.560428
75%	0.588857
max	1.000000

	net_profit_before_tax_to_paid_in_capital \
count	5522.000000
mean	0.182976
std	0.031848
min	0.000000
25%	0.169444
50%	0.178580
75%	0.191777
max	1.000000

	net_income_to_stockholders_equity	operating_profit_paid_in_capital \
count	5522.000000	5522.000000
mean	0.840298	0.109170
std	0.015906	0.028728
min	0.000000	0.000000
25%	0.840125	0.096221
50%	0.841196	0.104166
75%	0.842362	0.116023
max	1.000000	1.000000

	total_asset_turnover	total_expense_to_assets
count	5522.000000	5522.000000
mean	0.141240	0.029226
std	0.101618	0.027929
min	0.000000	0.000000
25%	0.076462	0.014526
50%	0.118441	0.022676
75%	0.176912	0.035854
max	1.000000	1.000000

[8 rows x 23 columns]

```
[12]: print(f"Total companies: {len(df):}")
```

Total companies: 5522

```
[13]: df[bankruptcy_col].value_counts()
```

```
[13]: bankrupt
0      5344
1       178
Name: count, dtype: int64
```

```
[14]: print(f'Bankruptcy Rate: {df[bankruptcy_col].mean() * 100:.2f}%')
```

Bankruptcy Rate: 3.22%

```
[15]: # 220 Bankrupt companies
      # 6599 non-Bankrupt companies
```

```
[16]: key_ratios = [
      'debt_ratio_percentage',
      'liability_to_equity',
      'current_liability_to_current_assets',
      'net_income_to_total_assets',
      'roa_before_interest_percent_after_tax',
      'working_capital_to_equity'
      ]

fig, axes = plt.subplots(2, 3, figsize=(18,10))
axes = axes.flatten()

for idx, ratio in enumerate(key_ratios):
    q1 = df[ratio].quantile(.05)
    q3 = df[ratio].quantile(.95)
    df_filtered = df[(df[ratio] >= q1) & (df[ratio] <= q3)]

    axes[idx].hist(df_filtered[df_filtered[bankruptcy_col]==0][ratio].dropna(),
                   bins=30, alpha=0.6, label='Non-Bankrupt', color='green',
    ↪edgecolor='black')
    axes[idx].hist(df_filtered[df_filtered[bankruptcy_col]==1][ratio].dropna(),
                   bins=30, alpha=0.6, label='Bankrupt', color='red',
    ↪edgecolor='black')


    axes[idx].set_xlabel(ratio.replace('_', ' ').title(), fontsize=10)
    axes[idx].set_ylabel('Frequency', fontsize=10)
    axes[idx].set_title(f'{ratio.replace("-", " ").title()}', fontsize=11,
    ↪fontweight='bold')
```

```

axes[idx].legend()
axes[idx].grid(alpha=0.3)

plt.suptitle('Financial Ratios: Distribution Comparison', fontsize=16,
             fontweight='bold', y=1.00)
plt.tight_layout()
plt.show()

```



```

[17]: comparison = df.groupby(bankruptcy_col)[key_ratios].mean().round(3)
comparison.index = ['non-Bankrupt', 'Bankrupt']
print("Mean Value Comparisons:")
print(comparison.T)

```

Mean Value Comparisons:

	non-Bankrupt	Bankrupt
debt_ratio_percentage	0.111	0.188
liability_to_equity	0.280	0.297
current_liability_to_current_assets	0.031	0.062
net_income_to_total_assets	0.810	0.741
roa_before_interest_percent_after_tax	0.562	0.461
working_capital_to_equity	0.736	0.724

```

[18]: important_ratios = [
        'debt_ratio_percentage',
        'liability_to_equity',
        'net_income_to_total_assets',
        'roa_before_interest_percent_after_tax'
    ]

```

```

]

fig, axes = plt.subplots(2, 2, figsize=(16,12))
axes = axes.flatten()

for idx, ratio in enumerate(important_ratios):
    q1 = df[ratio].quantile(0.05)
    q3 = df[ratio].quantile(0.95)
    df_filtered = df[(df[ratio] >= q1) & (df[ratio] <= q3)]

    box_data = [df_filtered[df_filtered[bankruptcy_col]==0][ratio].dropna(),
                 df_filtered[df_filtered[bankruptcy_col]==1][ratio].dropna()]

    bp = axes[idx].boxplot(box_data, labels=['Non-Bankrupt', 'Bankrupt'],
                           patch_artist=True, widths=0.6)

    colors = ['lightgreen', 'lightcoral']
    for patch, color in zip(bp['boxes'], colors):
        patch.set_facecolor(color)
        patch.set_alpha(0.7)
    axes[idx].set_ylabel(ratio.replace('_', ' ').title(), fontsize=11,
→fontweight='bold')
    axes[idx].set_title(f'{ratio.replace("_", " ").title()}',
                        fontsize=13, fontweight='bold', pad=15)
    axes[idx].grid(axis='y', alpha=0.3)

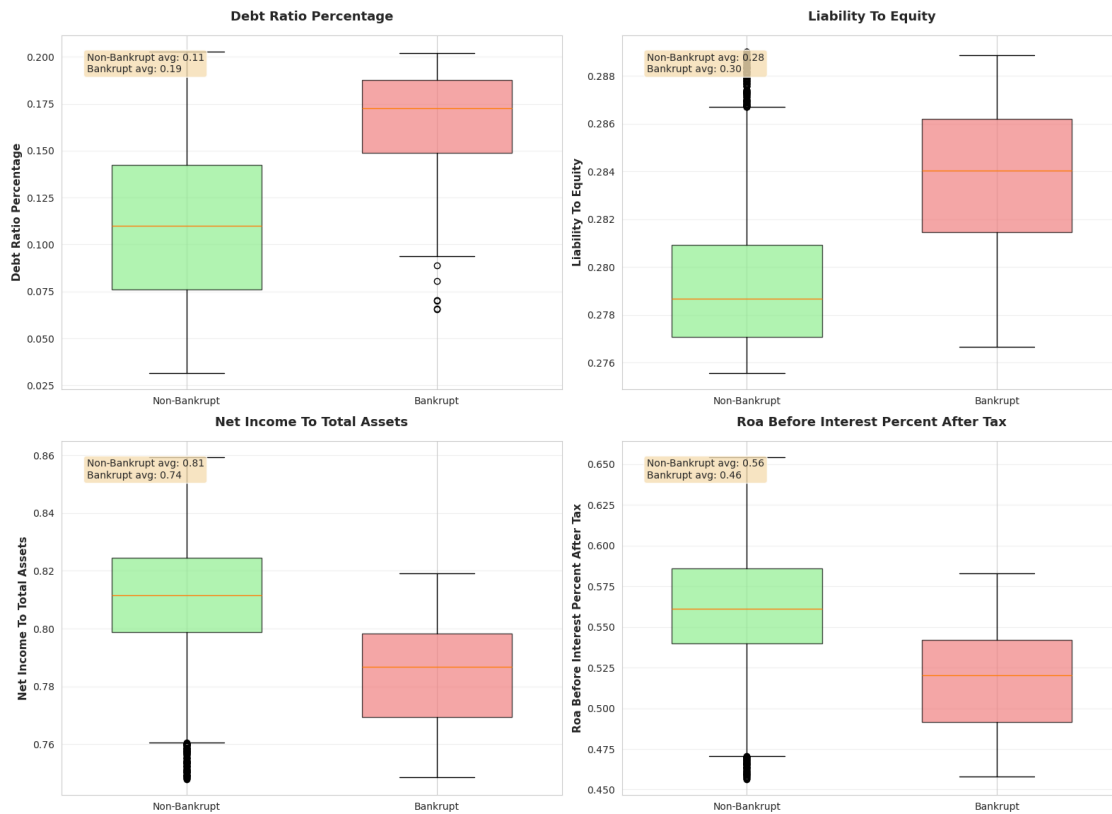
    mean_no_bank = df[df[bankruptcy_col]==0][ratio].mean()
    mean_bank = df[df[bankruptcy_col]==1][ratio].mean()

    textstr = f'Non-Bankrupt avg: {mean_no_bank:.2f}\nBankrupt avg: {mean_bank:.
→2f}'
    axes[idx].text(0.05, 0.95, textstr, transform=axes[idx].transAxes,
                   fontsize=10, verticalalignment='top',
                   bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.8))

plt.suptitle('Financial Ratios: Bankrupt vs Non-Bankrupt Companies',
             fontsize=16, fontweight='bold', y=1.00)
plt.tight_layout()
plt.show()

```


Financial Ratios: Bankrupt vs Non-Bankrupt Companies



```
[19]: print("Key Differences between Bankrupt vs. Non-Bankrupt:")
for ratio in important_ratios:
    mean_normal = df[df[bankruptcy_col]==0][ratio].mean()
    mean_bankrupt = df[df[bankruptcy_col]==1][ratio].mean()
    if mean_normal != 0:
        diff_pct = ((mean_bankrupt - mean_normal) / mean_normal * 100)
    else:
        0
    print(f'\n{ratio}: {diff_pct:+.1f}% difference')
```

Key Differences between Bankrupt vs. Non-Bankrupt:

debt_ratio_percentage: +70.2% difference

liability_to_equity: +6.2% difference

net_income_to_total_assets: -8.5% difference

roa_before_interest_percent_after_tax: -18.0% difference

```
[20]: numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
print(f'{len(numeric_cols)} numeric features')
```

23 numeric features

```
[21]: correlations = df[numeric_cols].corr()[bankruptcy_col].
    ↪ sort_values(ascending=False)
correlations = correlations.drop(bankruptcy_col)

print(correlations.head(10).to_string())
print("")
print(correlations.tail(10).to_string())
```

debt_ratio_percentage	0.251741
borrowing_dependency	0.214362
liability_to_equity	0.198140
current_liabilities_to_equity	0.191041
current_liability_to_current_assets	0.166677
equity_to_longterm_liability	0.151755
total_expense_to_assets	0.125832
total_asset_turnover	-0.068848
tax_rate	-0.101525
operating_gross_margin	-0.109776
working_capital_to_equity	-0.161235
per_share_net_profit_pre_tax	-0.194256
net_income_to_stockholders_equity	-0.195667
working_capital_to_total_assets	-0.199284
net_profit_before_tax_to_paid_in_capital	-0.199889
retained_earnings_to_total_assets	-0.206084
persistent_eps	-0.211888
net_worth_to_assets	-0.251741
roa_before_interest_percent_after_tax	-0.271845
net_income_to_total_assets	-0.300868

```
[22]: fig, axes = plt.subplots(1, 2, figsize=(16,8))

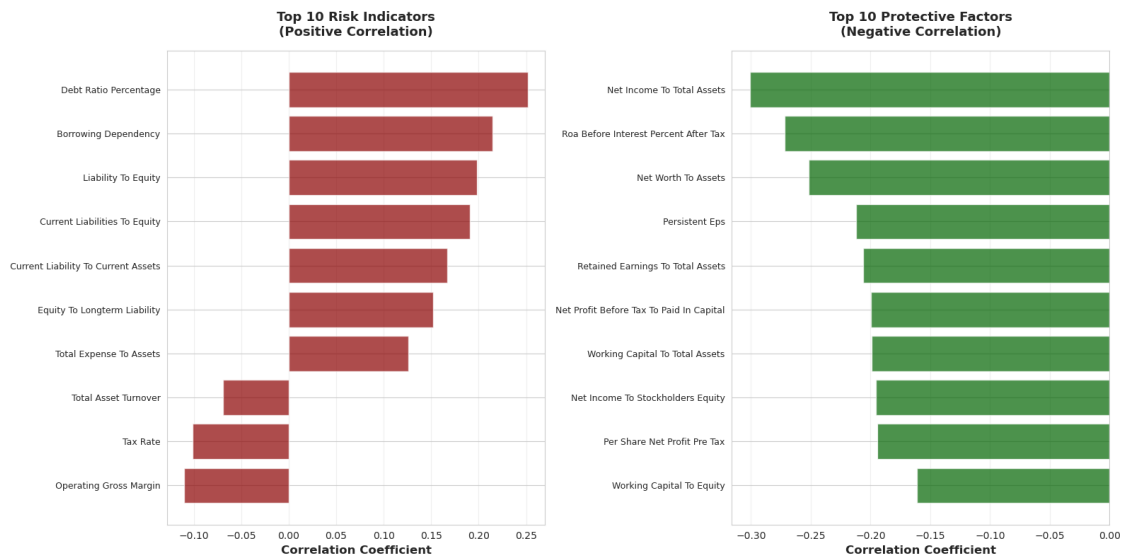
top_10_pos = correlations.head(10)
axes[0].barh(range(len(top_10_pos)), top_10_pos.values, color='darkred', alpha=0.
    ↪ 7)
axes[0].set_yticks(range(len(top_10_pos)))
axes[0].set_yticklabels([col.replace('_', ' ').title() for col in top_10_pos.
    ↪ index], fontsize=9)
axes[0].set_xlabel('Correlation Coefficient', fontsize=12, fontweight='bold')
axes[0].set_title('Top 10 Risk Indicators\n(Positive Correlation)',
    ↪ fontsize=13, fontweight='bold', pad=15)
axes[0].invert_yaxis()
axes[0].grid(axis='x', alpha=0.3)
```

```

# Top negative correlations (protective factors)
top_10_neg = correlations.tail(10).sort_values()
axes[1].barh(range(len(top_10_neg)), top_10_neg.values, color='darkgreen',
    ↳alpha=0.7)
axes[1].set_yticks(range(len(top_10_neg)))
axes[1].set_yticklabels([col.replace('_', ' ').title() for col in top_10_neg.
    ↳index], fontsize=9)
axes[1].set_xlabel('Correlation Coefficient', fontsize=12, fontweight='bold')
axes[1].set_title('Top 10 Protective Factors\n(Negative Correlation)',
    ↳fontsize=13, fontweight='bold', pad=15)
axes[1].invert_yaxis()
axes[1].grid(axis='x', alpha=0.3)

plt.tight_layout()
plt.show()

```



```

[23]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

feature_cols = [col for col in df.columns if col not in ['company_id',
    ↳bankruptcy_col]]
print(f'{len(feature_cols)} features for prediction')

X = df[feature_cols].copy()
y = df[bankruptcy_col].copy()
print(f'Target distribution: {y.value_counts().to_dict()}')

```

```

X = X.replace([np.inf, -np.inf], np.nan)
X = X.fillna(X.median())

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)
print(f"Training set: {X_train.shape[0]} samples ({y_train.mean()*100:.1f}% ↪bankruptcy)")
print(f"Test set: {X_test.shape[0]} samples ({y_test.mean()*100:.1f}% ↪bankruptcy)")

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

22 features for prediction
 Target distribution: {0: 5344, 1: 178}
 Training set: 4417 samples (3.2% bankruptcy)
 Test set: 1105 samples (3.3% bankruptcy)

```

[24]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import (classification_report, confusion_matrix,
                                   roc_auc_score, roc_curve, accuracy_score,
                                   precision_score, recall_score, f1_score)

      # Decided to use RF for this
      rf_model = RandomForestClassifier(
          n_estimators = 100,
          max_depth = 10,
          min_samples_split = 5,
          random_state = 42,
          class_weight = 'balanced',
          n_jobs = -1
      )

      rf_model.fit(X_train_scaled, y_train)

      y_pred = rf_model.predict(X_test_scaled)
      y_pred_proba = rf_model.predict_proba(X_test_scaled)[:,-1]

      # metrics
      accuracy = accuracy_score(y_test, y_pred)
      precision = precision_score(y_test, y_pred)
      recall = recall_score(y_test, y_pred)

```

```

f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_proba)

print(classification_report(y_test, y_pred, target_names=['non-Bankrupt',
↳ 'Bankrupt']))

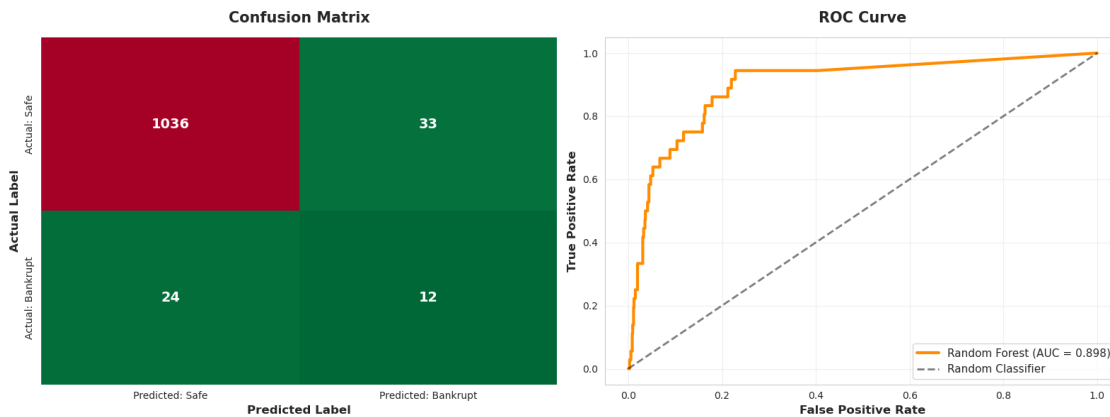
fig, axes = plt.subplots(1, 2, figsize=(16,6))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='RdYlGn_r', cbar=False, ax=axes[0],
            xticklabels = ['Predicted: Safe', 'Predicted: Bankrupt'],
            yticklabels = ['Actual: Safe', 'Actual: Bankrupt'],
            annot_kws = {"fontsize": 14, 'fontweight': 'bold'})
axes[0].set_title('Confusion Matrix', fontsize=15, fontweight='bold', pad=15)
axes[0].set_ylabel('Actual Label', fontsize=12, fontweight='bold')
axes[0].set_xlabel('Predicted Label', fontsize=12, fontweight='bold')

fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
axes[1].plot(fpr, tpr, linewidth=3, label=f'Random Forest (AUC = {roc_auc:.
↳ 3f})', color='darkorange')
axes[1].plot([0, 1], [0, 1], 'k--', linewidth=2, label='Random Classifier',
↳ alpha=0.5)
axes[1].set_xlabel('False Positive Rate', fontsize=12, fontweight='bold')
axes[1].set_ylabel('True Positive Rate', fontsize=12, fontweight='bold')
axes[1].set_title('ROC Curve', fontsize=15, fontweight='bold', pad=15)
axes[1].legend(fontsize=11, loc='lower right')
axes[1].grid(alpha=0.3)

plt.tight_layout()
plt.show()

```

	precision	recall	f1-score	support
non-Bankrupt	0.98	0.97	0.97	1069
Bankrupt	0.27	0.33	0.30	36
accuracy			0.95	1105
macro avg	0.62	0.65	0.63	1105
weighted avg	0.95	0.95	0.95	1105



```
[25]: feature_importance = pd.DataFrame({
    'feature': feature_cols,
    'importance': rf_model.feature_importances_
}).sort_values('importance', ascending=False)

print("Top 10 most important features:")
print(feature_importance.head(10).to_string(index=False))
```

Top 10 most important features:

feature	importance
borrowing_dependency	0.131195
persistent_eps	0.107448
net_income_to_total_assets	0.095930
per_share_net_profit_pre_tax	0.070879
retained_earnings_to_total_assets	0.062052
net_worth_to_assets	0.061965
debt_ratio_percentage	0.060598
liability_to_equity	0.048988
roa_before_interest_percent_after_tax	0.048194
net_profit_before_tax_to_paid_in_capital	0.047844

```
[26]: plt.figure(figsize=(12, 8))
top_features = feature_importance.head(10)

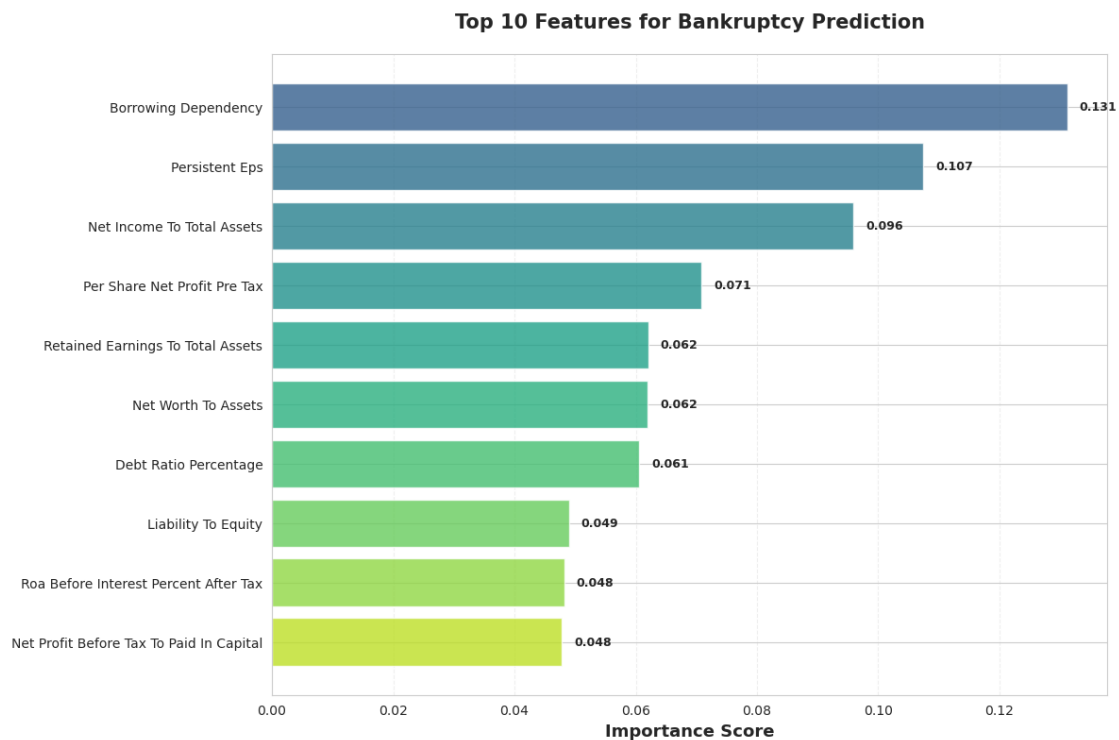
colors = plt.cm.viridis(np.linspace(0.3, 0.9, len(top_features)))
bars = plt.barh(range(len(top_features)), top_features['importance'],
    color=colors, alpha=0.8)

plt.yticks(range(len(top_features)), [f.replace('_', ' ').title() for f in
    top_features['feature']], fontsize=10)
plt.xlabel('Importance Score', fontsize=13, fontweight='bold')
plt.title('Top 10 Features for Bankruptcy Prediction', fontsize=15,
    fontweight='bold', pad=20)
```

```
plt.gca().invert_yaxis()
plt.grid(axis='x', alpha=0.3, linestyle='--')

# Add value labels
for i, (idx, row) in enumerate(top_features.iterrows()):
    plt.text(row['importance'] + 0.002, i, f"{row['importance']:.3f}",
            va='center', fontsize=9, fontweight='bold')

plt.tight_layout()
plt.show()
```



```
[27]: obj_predict = s3_client.get_object(
        Bucket=bucket_name,
        Key = 'machine-learning-data/company_profiles_to_predict_unlabeled.csv'
    )

df_predict = pd.read_csv(obj_predict['Body'])
df_predict.head()
```

```
[27]: borrowing_dependency    company company_id \
0          0.458819    western corp    id_6988
1          0.379304  design solutions    id_7413
2          0.384999    innocore        id_8801
```

3	0.374219	pharmasolve	id_9614
4	0.370253	ninetech	id_9131

	current_liabilities_to_equity	current_liability_to_current_assets	\
0	0.372218	0.060766	
1	0.333345	0.041220	
2	0.337392	0.060765	
3	0.329804	0.030201	
4	0.328093	0.021710	

	debt_ratio_percentage	equity_to_longterm_liability	liability_to_equity	\
0	0.269039	0.216878	0.337315	
1	0.161865	0.120812	0.282763	
2	0.216102	0.120561	0.292504	
3	0.108202	0.114508	0.278607	
4	0.058591	0.110933	0.276423	

	net_income_to_stockholders_equity	net_income_to_total_assets	...	\
0	0.765967	0.641804	...	
1	0.840533	0.800780	...	
2	0.829980	0.736619	...	
3	0.841459	0.815350	...	
4	0.840487	0.803647	...	

	operating_profit_per_share	per_share_net_profit_pre_tax	persistent_eps	\
0	0.050566	0.106274	0.154297	
1	0.107727	0.175868	0.217737	
2	0.099910	0.128945	0.161482	
3	0.104796	0.180462	0.225206	
4	0.095188	0.173157	0.218398	

	retained_earnings_to_total_assets	roa_before_interest_percent_after_tax	\
0	0.879445	0.344636	
1	0.933467	0.544320	
2	0.911441	0.445704	
3	0.935449	0.570922	
4	0.935200	0.545137	

	tax_rate	total_asset_turnover	total_expense_to_assets	\
0	0.000000	0.068966	0.079232	
1	0.564328	0.107946	0.038927	
2	0.000000	0.100450	0.092802	
3	0.157607	0.218891	0.025482	
4	0.005871	0.154423	0.029884	

	working_capital_to_equity	working_capital_to_total_assets
0	0.718867	0.720006

1	0.733645	0.769527
2	0.728731	0.729416
3	0.734946	0.797005
4	0.735182	0.813738

[5 rows x 24 columns]

```
[28]: if 'company_name' in df_predict.columns:
        company_names = df_predict['company_name'].copy()
    elif 'company' in df_predict.columns:
        company_names = df_predict['company'].copy()
    else:
        company_names = [f"Company_{i+1}" for i in range(len(df_predict))]

    if 'company_id' in df_predict.columns:
        company_ids = df_predict['company_id'].copy()
    else:
        company_ids = [f"ID_{i+1}" for i in range(len(df_predict))]

[29]: X_predict = df_predict[feature_cols].copy()

X_predict = X_predict.replace([np.inf, -np.inf], np.nan)
X_predict = X_predict.fillna(X_predict.median())

X_predict_scaled = scaler.transform(X_predict)

predictions = rf_model.predict(X_predict_scaled)
prediction_probabilities = rf_model.predict_proba(X_predict_scaled)[: ,1]

results = pd.DataFrame({
    'company_id': company_ids,
    'company_name': company_names,
    'bankruptcy_prediction': predictions,
    'bankruptcy_probability': prediction_probabilities,
    'bankruptcy_risk_pct': prediction_probabilities * 100
})

def categorize_risk(prob):
    if prob < .2:
        return 'LOW RISK'
    elif prob < .5:
        return 'MEDIUM RISK'
    else:
        return 'HIGH RISK'

results['risk_category'] = results['bankruptcy_probability'].
    ↪ apply(categorize_risk)
```

```

results = results.sort_values('bankruptcy_probability', ascending=False)

print("Bankruptcy Predictions")
print(results[['company_name', 'bankruptcy_risk_pct', 'risk_category']].
    ↳to_string(index=False))

```

Bankruptcy Predictions

company_name	bankruptcy_risk_pct	risk_category
western corp	87.746471	HIGH RISK
design solutions	68.107294	HIGH RISK
innocore	59.094088	HIGH RISK
Hallandall ag.	3.967326	LOW RISK
Highwood & Hart	1.922494	LOW RISK
songster inc	1.000000	LOW RISK
ninetech	0.000000	LOW RISK
pharmasolve	0.000000	LOW RISK
rogers and sons	0.000000	LOW RISK
Foster & Kruse	0.000000	LOW RISK

```

[30]: fig, axes = plt.subplots(1, 2, figsize=(16,6))

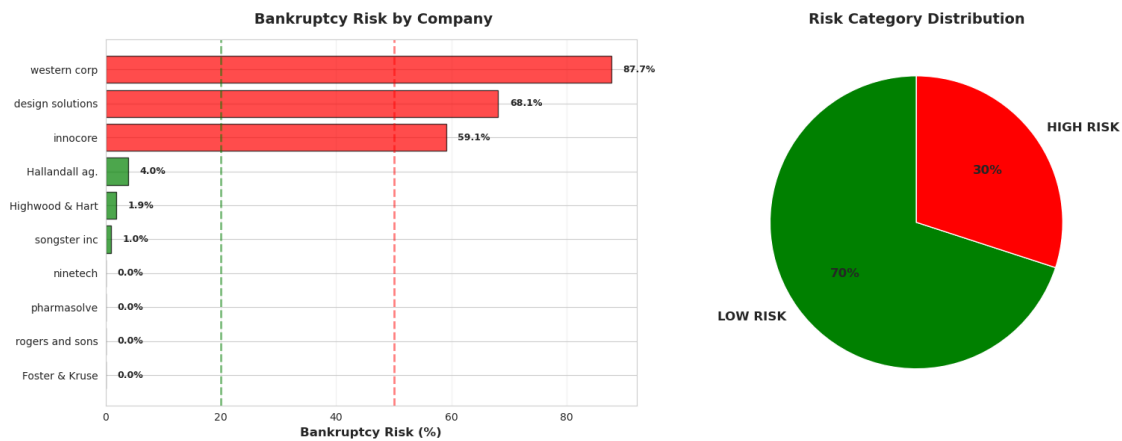
colors = ['red' if p > .5 else 'orange' if p > .2 else 'green' for p in
    ↳results['bankruptcy_probability']]
axes[0].barh(range(len(results)), results['bankruptcy_risk_pct'], color=colors,
    ↳alpha=0.7, edgecolor='black')
axes[0].set_yticks(range(len(results)))
axes[0].set_yticklabels(results['company_name'], fontsize=10)
axes[0].set_xlabel('Bankruptcy Risk (%)', fontsize=12, fontweight='bold')
axes[0].set_title(' Bankruptcy Risk by Company', fontsize=14, fontweight='bold',
    ↳pad=15)
axes[0].axvline(x=20, color='green', linestyle='--', linewidth=2, alpha=0.5)
axes[0].axvline(x=50, color='red', linestyle='--', linewidth=2, alpha=0.5)
axes[0].grid(axis='x', alpha=0.3)
axes[0].invert_yaxis()

for i, (idx, row) in enumerate(results.iterrows()):
    axes[0].text(row['bankruptcy_risk_pct'] + 2, i,
        ↳f"{row['bankruptcy_risk_pct']:.1f}%",
            va='center', fontsize=9, fontweight='bold')

risk_counts = results['risk_category'].value_counts()
colors_pie = ['green', 'red', 'orange']
axes[1].pie(risk_counts.values, labels=risk_counts.index, autopct='%1.0f%%',
    colors=colors_pie, startangle=90, textprops={'fontsize': 12,
    ↳'fontweight': 'bold'})
axes[1].set_title('Risk Category Distribution', fontsize=14, fontweight='bold',
    ↳pad=15)

```

```
plt.tight_layout()
plt.show()
```



```
[31]: low_risk = results[results['risk_category'] == 'LOW RISK']
medium_risk = results[results['risk_category'] == 'MEDIUM RISK']
high_risk = results[results['risk_category'] == 'HIGH RISK']

print("These companies show LOW bankruptcy risk:")
for idx, row in low_risk.iterrows():
    print(f"{row['company_name']}:")
    print(f" Bankruptcy Risk: {row['bankruptcy_risk_pct']:.1f}%")
```

These companies show LOW bankruptcy risk:

Hallandall ag.:

Bankruptcy Risk: 4.0%

Highwood & Hart:

Bankruptcy Risk: 1.9%

songster inc:

Bankruptcy Risk: 1.0%

ninetech:

Bankruptcy Risk: 0.0%

pharmasolve:

Bankruptcy Risk: 0.0%

rogers and sons:

Bankruptcy Risk: 0.0%

Foster & Kruse:

Bankruptcy Risk: 0.0%

```
[32]: print("These companies show MEDIUM bankruptcy risk:")
if len(medium_risk) > 0:
    for idx, row in medium_risk.iterrows():
```

```

        print(f"{row['company_name']}")
        print(f" Bankruptcy Risk: {row['bankruptcy_risk_pct']:.1f}%")
    else:
        print("No companies available")

```

These companies show MEDIUM bankruptcy risk:

No companies available

```

[33]: print("These companies show HIGH bankruptcy risk:")
      if len(high_risk) > 0:
          for idx, row in high_risk.iterrows():
              print(f"{row['company_name']}")
              print(f" Bankruptcy Risk: {row['bankruptcy_risk_pct']:.1f}%")

```

These companies show HIGH bankruptcy risk:

western corp

Bankruptcy Risk: 87.7%

design solutions

Bankruptcy Risk: 68.1%

innocore

Bankruptcy Risk: 59.1%

```

[34]: print(f"Total companies evaluated: {len(results)}")
      print(f"Recommended for investment: {len(low_risk)}")
      print(f"Not recommended: {len(high_risk)}")

```

Total companies evaluated: 10

Recommended for investment: 7

Not recommended: 3

```

[35]: summary_table = results[['company_name', 'bankruptcy_risk_pct',
    ↪ 'risk_category']].copy()
      summary_table['recommendation'] = summary_table['risk_category'].map({
          'LOW RISK' : 'INVEST',
          'MEDIUM RISK' : 'DO RESEARCH',
          'HIGH RISK' : 'AVOID'
      })
      print(summary_table.to_string(index=False))

```

company_name	bankruptcy_risk_pct	risk_category	recommendation
western corp	87.746471	HIGH RISK	AVOID
design solutions	68.107294	HIGH RISK	AVOID
innocore	59.094088	HIGH RISK	AVOID
Hallandall ag.	3.967326	LOW RISK	INVEST
Highwood & Hart	1.922494	LOW RISK	INVEST
songster inc	1.000000	LOW RISK	INVEST
ninetech	0.000000	LOW RISK	INVEST
pharmasolve	0.000000	LOW RISK	INVEST
rogers and sons	0.000000	LOW RISK	INVEST

Foster & Kruse	0.000000	LOW RISK	INVEST
----------------	----------	----------	--------

```
[36]: csv_buffer = StringIO()
results.to_csv(csv_buffer, index=False)

s3_client.put_object(
    Bucket = bucket_name,
    Key = 'predictions/bankruptcy_predictions.csv',
    Body = csv_buffer.getvalue()
)

print('predictions saved to s3')
```

predictions saved to s3

```
[37]: # Summary metrics for presentation (potential)
print(f"Model Accuracy: {accuracy*100:.1f}%")
print(f"ROC-AUC Score: {roc_auc:.3f}")
print(f"Companies to invest in: {len(low_risk)}")
print(f"Companies to avoid: {len(high_risk)}")
print(f"Top 3 risk indicators: {' '.join(feature_importance.head(3)['feature'].
    ↳tolist())}")
```

Model Accuracy: 94.8%

ROC-AUC Score: 0.898

Companies to invest in: 7

Companies to avoid: 3

Top 3 risk indicators: borrowing_dependency, persistent_eps,
net_income_to_total_assets